

# nullhat 2025: notThat

@elesquina

The jail reads one line of Python code and enforces a maximum length of 50 characters. It also rejects some substrings such as `import`, and it blocks characters like underscores and quotes. After the filter, it sets `sys.stdout` to a string, which makes normal output unusable, but exceptions still print to `stderr`. The process ends after one call to `chall()`, but inside the executed code we can call `chall()` again, so we can send many short inputs.

The useful observation is that global state survives between these nested calls. The module `string` is already imported at the top level, so it is a perfect scratchpad because we can attach attributes to it and reuse them later. This lets us build forbidden tokens without typing forbidden characters, for example `chr(95)` gives the underscore, and concatenation can create strings like `__builtins__` and `__import__`.

I split the exploit into multiple < 50-character inputs. Each input stores one intermediate value in `string.*`, then immediately calls `chall()` to get the next input evaluated in the same global context. At the end I use `glob` to find the randomized flag filename `flag-*` and I leak it by raising an assertion with the file contents as the message, because the traceback is printed even when `stdout` is broken.

The exact sequence of inputs is the following, with each line entered as one separate prompt:

```
string.u=chr(95);chall()
string.d=string.u*2;chall()
string.k=string.d+'builtins'+string.d;chall()
string.B=locals().get(string.k);chall()
string.m=string.d+('im'+port')+string.d;chall()
string.I=getattr(string.B,string.m);chall()
string.G=string.I('glob');chall()
string.F=next(string.G.iglob('flag-*'));chall()
assert 0,open(string.F).read()
```